AI-use disclosure

As permitted by Section 10 of the problem statement, an AI coding assistant was used during the hackathon window, for:

- help with implementation and with running experiments;
- drafting documentation.

Independent verification of all results

- All reported costs come from the unmodified official `scorer.py`. Its four published verification vectors are re-checked in `tests/test_pipeline.py`.
- The per-approach traffic decomposition is checked against the scorer on 1,240 random (scenario, schedule) pairs.
- The QUBO↔Ising conversion is tested. A separate test proves that the default QUBO is built without any joint-schedule cost evaluation.
- Every training simulation counts against the 300-evaluation budget; this is asserted by wrapping the simulator.
- The trained numpy state matches the exported Qiskit circuit ($|\Delta p| \le 4 \cdot 10^{-15}$), and this is asserted on every run.
- The OpenQASM 3 round trip is tested, and the counts are checked statistically against the circuit.
- Hyper-parameters were frozen on tuning scenarios before the held-out (seed 2026), ablation (seed 99), R and H runs.
- Every number in the report is regenerated by the commands in `README.md`.

Integrity

- No schedule, scenario-specific value or optimum is hard-coded.
- The QAOA objective uses QUBO energies only.
- The official cost is used solely for benchmarking and for the separately labelled hybrid result.